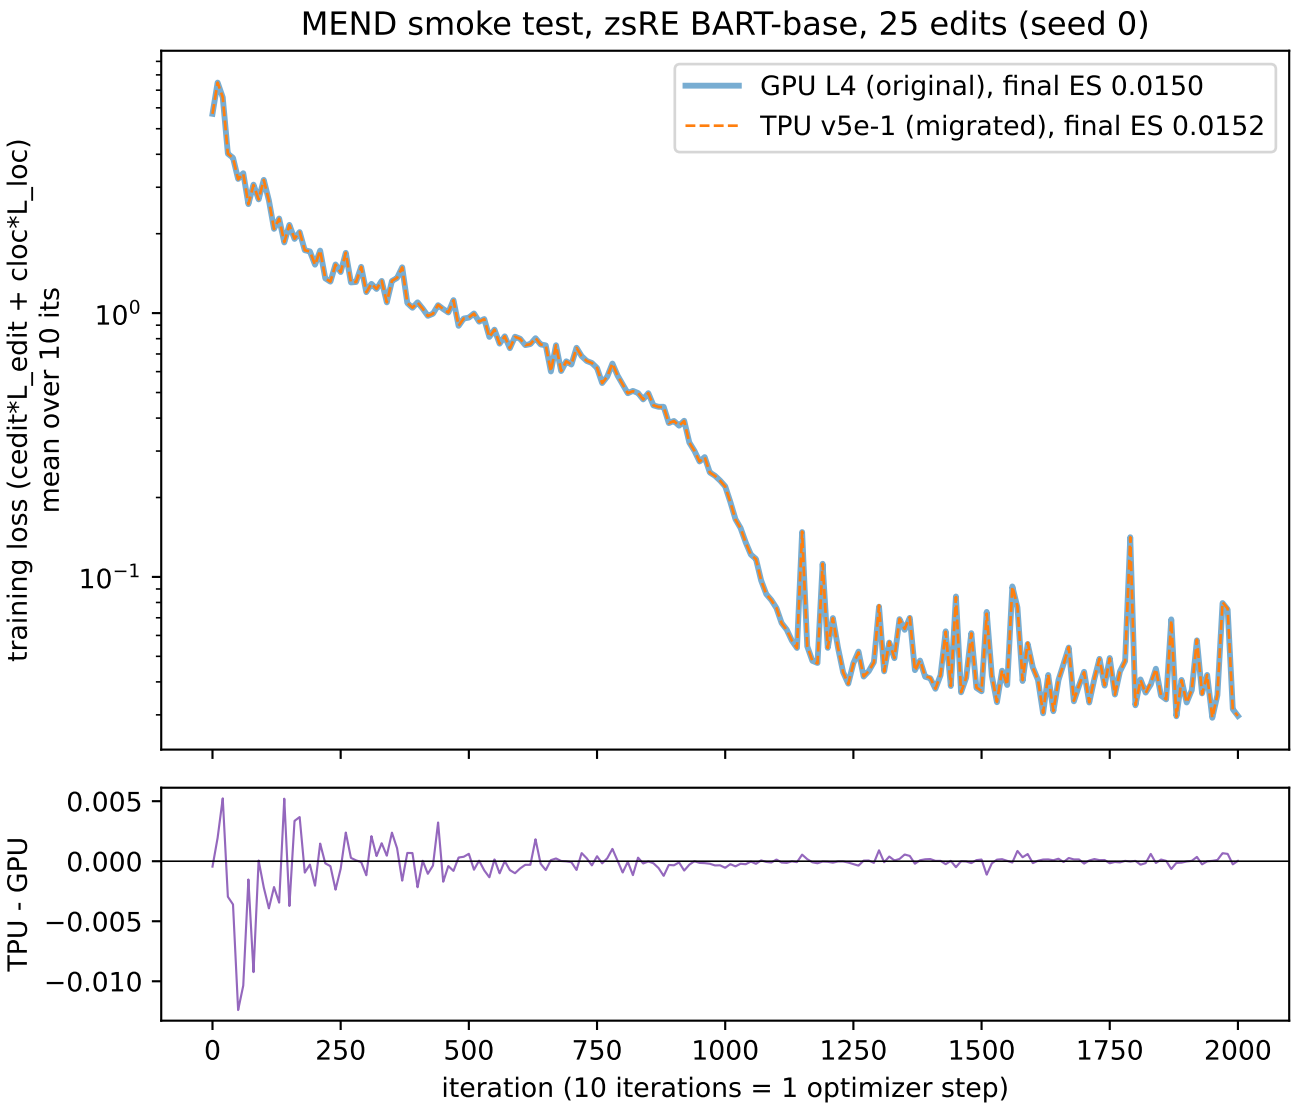
MEND smoke test, zsRE BART-base, 25 edits (seed 0)
GPU L4 (original), final ES 0.0150
TPU v5e-1 (migrated), final ES 0.0152
training loss (cedit*L_edit + cloc*L_loc)
mean over 10 its
10^{0}
10^{-1}
TPU - GPU
0.005
0.000
−0.005
−0.010
0
250
500
750
1000
1250
1500
1750
2000
iteration (10 iterations = 1 optimizer step)